

Intent Engineering for AI Agents: Landscape Survey (Early 2026)

Executive Summary

Intent engineering is emerging as a distinct discipline at the intersection of product management, software engineering, and AI agent design. Rather than telling agents *what to do* step by step, intent engineering focuses on specifying *what success looks like* — objectives, outcomes, constraints, and decision authority — so agents can reason autonomously within defined boundaries. As of February 2026, the landscape includes two dedicated platforms (Pathmode and Augment's Intent), a growing ecosystem of agent orchestration frameworks that handle intent concepts differently, two major interoperability protocols (MCP and A2A) under Linux Foundation governance, evaluation tools moving toward trace-based reasoning analysis, early academic formalization, and a small but vocal community of practitioners coalescing around the concept. No single standard has won, but the pieces are assembling rapidly.

Dedicated Intent Engineering Platforms

Pathmode

Pathmode positions itself as the first platform purpose-built for intent engineering, turning raw user feedback (support tickets, interview quotes, analytics) into structured specifications that AI coding agents can execute. The core innovation is the **IntentSpec** format — an open standard published at intentspec.org — that uses Markdown with YAML frontmatter and a formal JSON Schema for validation.[1][2]

The IntentSpec has five parts:[3]

1. **Objective** — The problem being solved and why it matters
2. **User Goal** — What the user wants to achieve

3. **Outcomes** — Measurable success states
4. **Edge Cases** — Boundary conditions and error states the implementation must handle
5. **Verification** — Specific test cases and acceptance criteria

IntentSpec is explicitly open and format-agnostic: specs can be written by hand in any text editor and consumed by any LLM or agent tool. Pathmode differentiates from [AGENTS.md](#) files by offering machine-readable structure with automated CI/CD validation, rather than unstructured text with no formal verification.[1]

In February 2026, Pathmode shipped the **Intent Compiler for Claude Code** — a zero-config MCP tool that turns Claude Code into a Socratic product thinking partner. It asks pointed questions, challenges vague language, and builds a structured intent spec through conversation. When the spec is ready, it saves an intent.md file and optionally exports as .cursorsrules or CLAUDE.md. The MCP server (npm [@pathmode/mcp-server@1.2.2](#)) works without an API key for the compiler tools; team features like dependency graphs, workspace strategy, and governance activate with an API key. Pathmode is listed on both the Official MCP Registry and the Cursor Directory.[4]

Pricing: Free tier available (\$0/month starting), with a subscription model for team features. IntentSpec itself is fully open with no vendor lock-in.[5][1]

Claude Code compatibility: Direct — the Intent Compiler is an MCP server designed specifically for Claude Code, and IntentSpec files work natively as [CLAUDE.md](#) context.[4][1]

Intent by Augment Code

Augment's **Intent** is a developer workspace designed as "what comes after the IDE" — a coordinated environment for orchestrating multiple AI agents. Launched in public beta in February 2026 for macOS (Windows waitlist), Intent reimagines development around spec-driven agent coordination rather than manual coding.[6]

The architecture works through a structured agent team:[7][6]

- A **Coordinator agent** uses Augment's Context Engine to understand the task and propose a plan as a spec
- **Implementor agents** execute work in parallel within isolated git worktrees
- A **Verifier agent** checks results against the spec before returning work for human review

The key differentiator is the **living spec** — a specification that evolves as agents make progress. Unlike static prompts that go stale, the spec is a shared document that agents read from and update, keeping all participants aligned as plans change. Full git workflow integration means developers can go from prompt to commit to PR to merge without leaving the app.[6]

Intent supports BYOA (Bring Your Own Agent): while optimized for Augment's own agent ("Auggie"), it also works with Claude Code, Codex, and OpenCode. Credits are charged at the same rate as the Augment CLI. Augment recommends installing their Context Engine MCP for developers using other agents.[7]

Claude Code compatibility: Supported as a BYOA provider. Claude Code sessions can run within Intent workspaces, benefiting from the coordinated spec and workspace isolation.[7]

Other Tools With Intent Positioning

OpenSpec is another spec-driven development architecture that addresses context loss in agent-driven workflows, though it has lower visibility than Pathmode or Intent.[8]

IBM's Objective-Validation Protocol is an emerging concept articulated by Ismael Faro of IBM Research, describing a shift "from vibe coding to Objective-Validation Protocol" where users define goals, agents execute autonomously, and humans approve at critical checkpoints. IBM frames this as the foundation for an "Agentic Operating System (AOS)" that would standardize orchestration, safety, compliance, and resource governance across agent swarms. This remains a conceptual framework rather than a shipping product as of February 2026.[9][10]

Agent Orchestration Frameworks

Each major framework handles the relationship between human intent and agent execution differently. The following comparison captures how each approaches intent specification as of early 2026.

Fram ewor k	Intent Approach	Orchestratio n Model	Claude Code Compatibl e
Lang Grap h	Graph-based state machines; intent encoded in node/edge structure and state transitions	Directed graph with conditional branching[11]	Via LangChain integrations
Crew AI	Role-based; intent expressed through agent role definitions and task descriptions	Role assignment with Flows for workflow control[11]	Via tool integrations
Googl e ADK	Hierarchical delegation; intent in agent role descriptions and routing logic	Multi-agent hierarchy with LLM-driven routing[12]	Via LiteLLM adapter
Claud e Agen t SDK	System prompts and tool permissions; intent as conversational context	Single agent with subagent spawning[13][14]	Native — wraps Claude Code CLI
Agen t Team s	Shared task list and messaging; intent set by team lead, refined through collaboration	Full mesh multi-agent with team lead[15][16]	Native — experimental Claude Code feature

LangGraph

LangGraph models agent workflows as directed graphs where nodes represent processing steps and edges define transitions. Intent is essentially structural: the graph topology itself encodes what should happen, in what order, and under what conditions. This provides maximum control — developers can define human-in-the-loop pause/resume hooks within workflow graphs — but requires the intent to be pre-decomposed into graph logic. LangGraph excels at complex decision-making pipelines with conditional branches and is the strongest choice for enterprise applications requiring strict state management and audit trails.[11][17]

CrewAI

CrewAI takes an organizational metaphor: agents are assigned roles with specific responsibilities, mimicking team structures. Intent is expressed through role definitions ("You are a researcher responsible for...") and task descriptions, making it intuitive for PMs to think about. CrewAI's Flows feature adds event-driven pipeline control alongside autonomous Crews, combining structured orchestration with flexible agent reasoning. The trade-off is less fine-grained control over state transitions compared to LangGraph.[11]

Google Agent Development Kit (ADK)

Google's ADK, open-sourced at Cloud NEXT 2025, enables multi-agent hierarchies where agents delegate tasks based on context and capability. Intent is encoded through agent role descriptions and routing logic — a root agent can dynamically route requests to specialist sub-agents based on the content of the request. ADK supports flexible orchestration combining deterministic workflow agents with LLM-driven routing, and provides native support for the A2A protocol for inter-agent communication. Available in Python, Go, and Java.[12][18]

Anthropic's Claude Agent SDK

The Claude Agent SDK is Anthropic's official Python toolkit for building agents on top of Claude Code. Released alongside Claude Sonnet 4.5, it provides async streaming interfaces, custom tools via in-process MCP servers, a hook system for injecting logic before/after tool execution, and fine-grained permission control. Intent is specified through system prompts, tool availability, and permission boundaries — essentially, the same infrastructure that powers Claude Code is exposed for custom agent development. The SDK wraps the Claude Code CLI as a subprocess, with the CLI bundled in the package.[13][14]

Claude Code Agent Teams

Agent Teams is an experimental feature that coordinates multiple Claude Code sessions working together. One session acts as the **team lead** (understanding the big picture, delegating tasks, synthesizing results), while multiple **teammate sessions** operate as full independent Claude Code instances with their own context windows. Intent flows through a shared task list and direct mailbox system — any teammate can message any other teammate without routing through the lead.[19][15][16]

The communication model is full mesh rather than hub-and-spoke, making Agent Teams best suited for collaborative exploration, research, debugging with competing hypotheses, and cross-layer coordination. Enabling the feature requires setting `experimental.agentTeams: true` in Claude Code settings. Token costs are higher than subagent approaches since each teammate is a separate Claude instance.[15][16][19]

Standards and Protocols

Model Context Protocol (MCP)

MCP, introduced by Anthropic in November 2024, has rapidly become the universal standard for connecting AI models to tools, data, and applications. With over 10,000 published MCP servers and adoption across Claude, Cursor, VS Code, ChatGPT, Gemini, and Microsoft

Copilot, MCP provides the connective tissue that makes intent engineering practical.[20]

In December 2025, Anthropic donated MCP to the **Agentic AI Foundation (AAIF)** under the Linux Foundation, alongside Block's Goose and OpenAI's [Agents.md](#). The protocol's specification (latest: 2025-06-18) defines servers that expose Resources (context/data), Prompts (templated workflows), Tools (executable functions), Sampling (server-initiated LLM interactions), Roots (filesystem boundaries), and Elicitation (user information requests).[21][22][20]

For intent engineering, MCP serves as the **delivery mechanism** rather than the specification layer. Pathmode's IntentSpec, for example, can be delivered to agents as MCP context. MCP enables intent but does not define it — there is currently no intent-specific schema within the MCP specification itself.[3]

Agent-to-Agent Protocol (A2A)

Google's A2A protocol, announced in April 2025 and contributed to the Linux Foundation in June 2025, addresses inter-agent communication — how agents discover each other's capabilities, exchange information, and coordinate actions. Over 100 technology partners support A2A, including Atlassian, Salesforce, ServiceNow, and LangChain.[23][24]

A2A uses **Agent Cards** (JSON format) for capability discovery, supports task management with defined lifecycle states, and enables agent-to-agent collaboration via context and instruction sharing. It's designed to complement MCP: where MCP connects agents to tools, A2A connects agents to each other. The protocol is built on HTTP, SSE, and JSON-RPC.[25][24]

However, adoption has been slower than MCP's rapid uptake. IBM's Agent Communication Protocol (ACP), launched in March 2025, has since merged with A2A under the Linux Foundation, consolidating the agent-to-agent protocol space.[26][27][28]

Emerging Intent Standards

There is no widely adopted standard specifically for intent specification as of February 2026. The closest candidates are:

- **IntentSpec** (Pathmode) — open, JSON Schema-validated, Markdown-based; the most formalized intent specification format available[1]
- **IBM's Objective-Validation Protocol** — conceptual framework for goal-setting and checkpoint validation[10]
- **Agents.md** (OpenAI, now in AAIF) — unstructured agent context files that carry intent but lack formal validation[22]

The gap between MCP (tool connectivity) and A2A (agent communication) leaves intent specification as an unfilled layer in the protocol stack. Whether IntentSpec or a competing format fills this gap remains an open question.

Evaluation and Observability

Trace-Based Evaluation

The shift from evaluating agent outputs to evaluating agent reasoning is one of the most significant developments in the observability space. Leading platforms now capture full decision traces — every prompt, model output, tool call, and intermediate reasoning step — enabling developers to understand not just *what* an agent did but *why*. [29]

Platform	Approach	Best For
Braintrust	Evaluation-first with comprehensive traces; nested spans across multi-agent workflows; automated scoring[30]	Teams needing integrated eval + observability
LangSmith	Step-by-step inspection of decision paths; run replay and comparison; tight LangChain integration[31]	LangChain-based agent debugging
Arize AX	SDK-based "decision-level" visibility showing how internal state changes between tool calls[32]	Mission-critical agents requiring deep reasoning visibility
Langfuse	OpenTelemetry-compatible traces with cost/latency tracking; sequential chain visualization[31]	Teams optimizing performance and spend

These tools collectively enable **intent alignment monitoring** — checking whether an agent's reasoning chain stays aligned with its specified objectives. Braintrust, for example, supports configurable evaluation scores that can trigger alerts when quality degrades, enabling teams to catch regressions before users see them.[30]

Amazon's LLM Simulator for Intent Detection

Amazon developed an LLM simulator for evaluating customer service agent intent detection accuracy. The system uses LLM-driven virtual customer personas to simulate diverse scenarios and interactions, then compares the agent's detected intent against ground truth intents from anonymized historical customer interactions.[33]

The evaluation covers multiple layers:[33]

- **Intent correctness** — whether the orchestration agent understands user intents correctly

- **Task completion** — whether the agent's final response resolves the customer's issue
- **Multi-turn coherence** — topic adherence classification and topic adherence refusal metrics
- **Tool use accuracy** — correct tool selection, parameter accuracy, and multi-turn function calling sequence

This approach represents a broader industry shift toward evaluating agents as complete systems rather than isolated models, with intent detection accuracy as a first-class metric.

Instruction Adherence as a Metric

Salesforce predicts that instruction adherence will become a key reliability metric in 2026, with enterprises demanding "probabilistic adherence scores — categorized as high, low, or uncertain — to enable developers to refine their instructions accordingly". This framing treats intent specification as something that can be quantitatively measured against agent behavior, moving intent engineering from an art toward a measurable discipline.[34]

Academic Research

"Human Delegation Behavior in Human-AI Collaboration" (arXiv:2401.04729)

This paper by Spitzer et al. (January 2024, revised January 2025) is frequently cited in the intent engineering community — notably by Paweł Huryn — as evidence that providing contextual information significantly improves human-AI team performance. The study found that access to contextual information about both the AI's capabilities and the task domain changes delegation behavior and improves collaborative outcomes. While the paper focuses on human delegation rather than agent specification per se, its finding that *context beyond raw task specifications improves autonomous performance* has become a foundational reference for intent engineering practitioners.[35][36][37]

"Levels of Autonomy for AI Agents" (Columbia / Knight Institute)

This working paper defines five levels of escalating agent autonomy, characterized by the roles a user takes: **Operator, Collaborator, Consultant, Approver, and Observer**. The key argument is that autonomy should be treated as a **deliberate design decision**, separate from agent capability — a capable agent can still be designed to operate at low autonomy if it must consult its user before each action.[38]

The paper proposes **autonomy certificates** as a governance mechanism: third-party-issued certificates that communicate behavioral characteristics of an agent to other developers and agents. This framework directly supports intent engineering by providing a vocabulary for specifying how much decision authority an agent should have — a critical component of any intent specification.[38]

"Agentic Software Engineering" Textbook

Available as a free PDF at agenticse-book.github.io, this book (February 2026) defines Agentic Software Engineering as "the discipline of producing high-quality, reliable, trustworthy software from stochastic contributors, both AI and human". The thesis is explicit: "Success belongs to teams who set clear intent, manage risk boundaries, and demand evidence — not just those who type fastest". [39][40]

The book frames itself as "the missing manual that did not come with your agentic coding tools, including Claude Code" and addresses the four pillars of engineering (people, process, tools, and artifacts) rather than just coding techniques. It emphasizes that intent specification is necessary precisely because AI teammates are stochastic — they fail with some probability, requiring constraint-based engineering to achieve trusted reliability.[39]

Intent-Based System Design and Operation (arXiv:2502.05984)

This February 2025 paper proposes intent as a new abstraction within the context of system design and operation, encoding functional and operational requirements at a high level to automate design, implementation, operation, and evolution of systems. While broader than AI agent intent engineering, it formalizes many of the same concepts — using intent as the interface between human goals and system behavior.[41]

Community and Discourse

Key Voices

Paweł Huryn is the most systematic voice on intent engineering for AI agents. His Product Compass newsletter (119K+ subscribers) published "The Intent Engineering Framework for AI Agents" in January 2026, defining a seven-element structure:[36][42]

1. Objective — the problem to solve and why it matters
2. Desired Outcomes — observable state changes indicating success
3. Health Metrics — what must not degrade while pursuing outcomes
4. Strategic Context — the system being operated in
5. Constraints — steering (prompt-layer) vs. hard (orchestration-enforced)
6. Decision Types and Autonomy — which decisions require escalation
7. Stop Rules — when to halt, escalate, or complete

Huryn's formulation bridges product management (OKRs, empowered teams) with agent engineering, drawing explicitly on arXiv:2401.04729 and concepts from military "Auftragstaktik" (commander's intent). His framing — "context without intent is noise" — has become a touchstone for the community. He emphasizes that intent engineering is not about more detailed instructions, but about giving agents the reasoning framework to handle unknowns.[43][36]

Nate B Jones runs a Substack newsletter ("Nate's Substack" / "AI News & Strategy Daily") focused on AI strategy and practical implementation. His content covers agent UI surfaces, intent-state-outcome loops, and predictions around memory breakthroughs and long-running agents. Jones emphasizes execution over tooling and has published guides on agent architecture, multi-agent orchestration, and context engineering.[44][45][46][47][48][49]

The Agentic SE Book authors ([agenticsse-book.github.io](https://github.com/agenticsse-book)) contribute the academic engineering perspective, arguing that intent clarity is the foundation of trustworthy agentic development.[40][39]

IBM researchers (particularly Ismael Faro and Maryam Ashoori) are shaping the enterprise perspective with concepts like the Objective-Validation Protocol and predictions about 2026 as "the year of operating AI agents at scale".[50][10]

Newsletters and Communities

Resource	Focus	Author/Or g
The Product Compass	Intent engineering frameworks, context engineering, AI PM	Paweł Huryn[42]
Nate's Substack	AI strategy, agent architecture, practical playbooks	Nate B Jones[49]
r/AugmentCodeAI	Agent orchestration, Intent workspace discussion	Reddit community [7]
ClaudeFast blog	Claude Code Agent Teams guides, team orchestration patterns	Community guides[15]
IBM Think / IBM Research Blog	Super agents, ACP/A2A protocol development, enterprise agentic AI	IBM[10] [51]

Forming Reading List

A canonical reading list is not yet established, but the emerging syllabus includes:

- Huryn's "Intent Engineering Framework for AI Agents" and "Context Engineering" posts[35][36]
 - arXiv:2401.04729 on contextual information in human-AI delegation[37]
 - "Levels of Autonomy for AI Agents" (Columbia/Knight)[38]
 - The Agentic Software Engineering textbook[39]
 - IntentSpec documentation at intentspec.org[1]
 - Anthropic's MCP specification[21]
 - Amazon's agent evaluation blog on intent detection with LLM simulators[33]
 - arXiv:2502.05984 on intent-based system design[41]
-

Landscape Assessment

What Exists Today

The intent engineering landscape in early 2026 is fragmented but directional. Two dedicated platforms (Pathmode and Augment Intent) have shipped products, one open standard (IntentSpec) is available with CI/CD validation, and one conceptual framework (Huryn's 7-element structure) provides a practitioner vocabulary. The protocol stack (MCP for tools, A2A for agents) is maturing under Linux Foundation governance but lacks an intent-specific layer.

What's Missing

- **No universal intent schema** — IntentSpec is the closest, but adoption beyond Pathmode's ecosystem is nascent
- **No runtime intent enforcement** — most intent specifications operate at the prompt layer; hard enforcement remains a custom engineering problem
- **No cross-framework intent portability** — an IntentSpec written for Claude Code doesn't automatically work as a CrewAI crew definition or a LangGraph state machine

- **Limited evaluation tooling for intent alignment** — trace-based tools monitor reasoning but don't systematically compare agent behavior against intent specifications
- **Sparse academic formalization** — the concept is referenced in industry but has no dedicated peer-reviewed research program

Trajectory

The convergence of several trends suggests intent engineering will solidify rapidly: the shift from "vibe coding" to structured development (IBM's framing), the maturation of agent orchestration frameworks that need better specification inputs, the growing observability stack that can measure intent adherence, and the professional community forming around Huryn's framework and Pathmode's tools. For anyone building with Claude Code today, the practical action is to adopt IntentSpec files (intent.md) as project artifacts alongside CLAUDE.md, and to treat intent specification as a first-class product skill rather than an afterthought.[34][10]

References

1. [IntentSpec | The Open Standard for Spec-Driven Development](#) - Do I need Pathmode to use IntentSpec? No. IntentSpec is an open standard. You can write [intent.md](#) fi...
2. [Pathmode: Know Exactly What to Build](#) - From raw feedback to specs your AI agents can actually ship. Pathmode turns user friction into build...
3. [IntentSpec — Definition & Guide | Pathmode Glossary](#) - The IntentSpec is a five-part structured specification format that captures product intent precisely...
4. [Changelog | Pathmode](#) - Connect Linear to link issues and projects with Intent Specs. Track how product decisions flow from ...
5. [Pathmode 2026 Pricing, Features, Reviews & Alternatives | GetApp](#) - Pathmode is a product discovery software that transforms unstructured user feedback into execution-r...
6. [Intent: A workspace for agent orchestration - Augment Code](#) - TL;DR: Intent is our vision for what comes after the IDE. It's a developer workspace designed for or...

7. [Introducing Intent: A workspace for agent orchestration - Reddit](#) - Intent is our vision for what comes after the IDE. It's a developer workspace designed for orchestra...
8. [OpenSpec Deep Dive: Spec-Driven Development Architecture ...](#) - Explore OpenSpec's core architecture, workflows, and enterprise applications to understand how Spec...
9. [Getting Super-Agents This Year, Says IBM Roundup - Forbes](#) - IBM predicts 2026 super-agents, control planes, and objective-validation protocols guiding autonomou...
10. [The trends that will shape AI and tech in 2026 | IBM](#) - Reporter Anabelle Nicoud spoke to several experts across AI, security, quantum and beyond to better ...
11. [CrewAI vs LangGraph vs AutoGen: Choosing the Right Multi-Agent ...](#) - In this tutorial, I will walk you through a detailed comparison of three leading multi-agent AI fram...
12. [Google Introduces Agent Development Kit \(ADK\) at Cloud NEXT 2025](#) - It supports: Multi-Agent by Design: Developers can define hierarchies where agents delegate tasks in...
13. [anthropics/claude-agent-sdk-python - GitHub](#) - Python SDK for Claude Agent. See the Claude Agent SDK documentation for more information. Installati...
14. [Introducing Claude Sonnet 4.5 - Anthropic](#) - Claude Sonnet 4.5 is the best coding model in the world, strongest model for building complex agents...
15. [Claude Code Agent Teams: The Complete Guide 2026](#) - Note on terminology: This post covers Claude Code's native Agent Teams feature, an experimental buil...
16. [Orchestrate teams of Claude Code sessions](#) - Agent teams let you coordinate multiple Claude Code instances working together. One session acts as ...
17. [The Best Open Source Frameworks For Building AI Agents in 2026](#) - Seven open source frameworks compared: LangGraph, OpenAI Agents SDK, AutoGen, CrewAI, Google ADK, Di...
18. [Google Launches Agent Development Kit for Go - InfoQ](#) - The Go ADK is an open-source toolkit that enables developers to build modular multi-agent systems in...
19. [How Multi-Agent Coding Actually Works - Kumar Gauraw](#) - Claude Code Agent Teams Explained: How Multi-Agent Coding Actually Works · Claude Code vs Codex in 2...

20. [Linux Foundation Announces the Formation of the Agentic AI ...](#) - The AAIF provides a neutral, open foundation to ensure agentic AI evolves transparently and collabor...
21. [Specification - What is the Model Context Protocol \(MCP\)?](#) - Model Context Protocol (MCP) is an open protocol that enables seamless integration between LLM appli...
22. [Why the Agentic AI Foundation \(AAIF\) Changes Everything for MCP](#) - The Linux Foundation's new Agentic AI Foundation (AAIF) brings MCP, Goose from Block, and [Agents.md](#) ...
23. [Linux Foundation Launches the Agent2Agent Protocol Project to ...](#) - The Agent2Agent protocol enables agentic AI interoperability and trusted agent communication across ...
24. [Announcing the Agent2Agent Protocol \(A2A\)](#) - The A2A protocol will allow AI agents to communicate with each other, securely exchange information,...
25. [Best of 2025: Google Cloud Unveils Agent2Agent Protocol](#) - Google Cloud's new Agent2Agent Protocol enables AI agents from different vendors to communicate seam...
26. [MCP and A2A: A Network Engineer's Mental Model for Agentic AI](#) - In April 2025, Google introduced A2A, providing a new protocol that allows agents to discover each o...
27. [Armand Ruiz's Post - LinkedIn](#) - In March 2025, my colleagues from IBM Research launched the Agent Communication Protocol (ACP), our ...
28. [ACP Joins Forces with A2A - Linux Foundation AI & Data](#) - IBM Research launched the Agent Communication Protocol (ACP) in March 2025 to power its BeeAI Platfo...
29. [Top 5 AI Agent Observability Platforms 2026 Guide | Articles | o-mega](#) - 5. LangSmith – Developer-Centric Agent Tracing by LangChain. LangSmith is the observability and eval...
30. [5 best AI agent observability tools for agent reliability in 2026 - Articles](#) - Braintrust - Best overall AI agent observability platform. Evaluation-first architecture with compre...
31. [15 AI Agent Observability Tools in 2026: AgentOps & Langfuse](#) - LangSmith captures full reasoning traces for LangChain-based agents, including prompts, retrieved co...
32. [Best AI Observability Tools for Autonomous Agents in 2026 - Arize AI](#) - Compare the best platforms for monitoring autonomous agents in 2026. Learn how to choose between pro...

33. [Evaluating AI agents: Real-world lessons from building agentic ...](#)
- To evaluate the agent's reasoning capability for intent detection, the Amazon team developed an LLM ...
34. [The Future of AI Agents: Top Predictions and Trends to Watch in 2026](#) - 2025 delivered enterprise AI that moved beyond simple prompts and reactive text generation into a ne...
35. [A Guide to Context Engineering for PMs - The Product Compass](#) - ... 2401.04729 that demonstrates how supplying strategic context beyond raw task specifications impr...
36. [The Intent Engineering Framework for AI Agents](#) - The Intent Engineering Framework shows how to structure objectives, outcomes, health metrics, constr...
37. [\[2401.04729\] Human Delegation Behavior in Human-AI Collaboration](#) - Abstract page for arXiv paper 2401.04729: Human Delegation Behavior in Human-AI Collaboration: The E...
38. [Levels of Autonomy for AI Agents Working Paper - arXiv](#) - We define five levels of escalating agent autonomy, characterized by the roles a user can take when ...
39. [\[PDF\] Agentic Software Engineering](#) - So this book is about Agentic Software Engineering, not agentic coding. It is about how to run softw...
40. [Agentic Software Engineering - The Future of Code](#) - A comprehensive guide to Agentic Software Engineering. Learn how AI agents are transforming the soft...
41. [Intent-based System Design and Operation](#) - ...industry. Yet, these systems remain complicated to design, build, operate, and improve. All these...
42. [The Product Compass | Paweł Huryn | Substack](#) - The #1 AI & PM newsletter. Actionable tips, resources, and step-by-step playbooks to learn by doing....
43. [Paweł Huryn's Post - LinkedIn](#) - Intent really is the compass. Without it, even good AI tends to drift. With clear intent, autonomy b...
44. [Nate B. Jones' AI Predictions for 2026: How I'm - Araptus Highlander](#) - Discover how Nate B. Jones' groundbreaking AI predictions are shaping how I retool every client with...
45. [Nate B Jones on AI Chaos and Unpredictability | Marc Wilson posted ...](#) - His Substack newsletter, including pieces like '2026 Sneak ... In videos such as 'The AI Moments Tha...

46. [Agents Will Kill Your UI by 2026--Unless You Build This Instead](#) - Where designers, PMs, and engineers must reskill around intent. For B2B SaaS leaders ... AI News & S...
47. [Executive Briefing: Your 2025 AI Agent Playbook in 10 Minutes ...](#) - This analysis is based on Nate B Jones's enterprise AI agent architecture framework, which argues th...
48. [Sitemap - 2025 - Nate's Substack](#) - Grab my Zero to Hero AI Agent Guide for NON-Tech: 28 Prompts, 4 Get Started Scenarios, Detailed Guid...
49. [Nate's Substack | Substack](#) - Daily newsletters from the cutting edge of AI, including breaking news, strategy, product breakdowns...
50. [IBM's Bold AI Predictions for 2026 and Beyond - The AI Innovator](#) - In a series of 2026 predictions shared with The AI Innovator, IBM executives argue that the industry...
51. [Introducing multiagent BeeAI - IBM Research](#) - IBM's ACP takes things a step further by introducing the ability to discover and run agents. ACP is ...